

Blockchains & Distributed Ledgers


Lecture 05

Aggelos Kiayias




Slide credits: AK, Dimitris Karakostas, Dionysis Zindros, Christos Nasikas

The Byzantine Generals Problem









[LSP82] Leslie Lamport, Robert Shostak, and Marshall Pease.
The byzantine generals problem

The Byzantine Generals Problem



The Byzantine Generals Problem




The Consensus Problem

Motivation for the Consensus Layer, I

- A transaction history and/or state of the service needs to be **agreed** by all servers.
- Servers may be operated by participants with **diverging interests**, in terms of the history of transactions and/or state of the service.

Motivation for the Consensus Layer, II

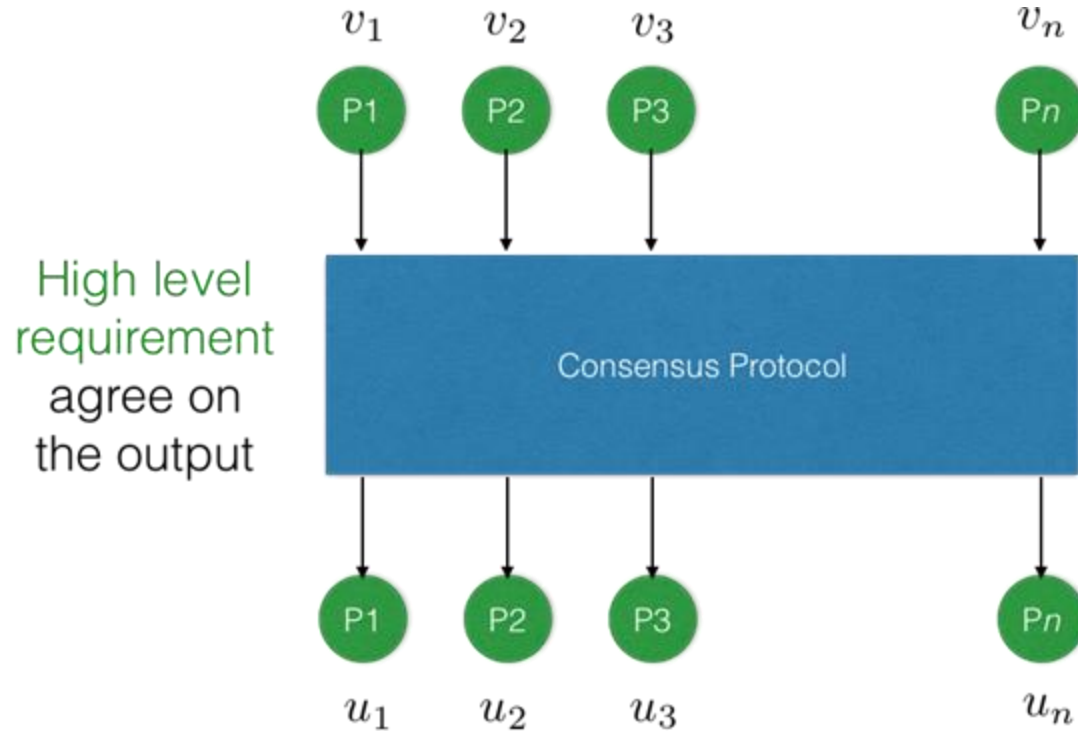


Consensus : Problem Statement

- A number (t) of the participating entities can diverge from the protocol.
- This has been called **Byzantine behaviour** in the literature.
- The properties of the protocol are defined in the presence of this “malicious” coalition of parties that attempts to disrupt the process for the “honest” parties.

$$H, |H| = n - t$$

The consensus problem



Study initiated by Lamport, Pease, Shostak 1982

Consensus Properties

- Termination $\forall i \in H(u_i \text{ is defined})$

Consensus Properties

- Termination $\forall i \in H (u_i \text{ is defined})$
- Agreement $\forall i, j \in H (u_i = u_j)$


Consensus Properties

- Termination $\forall i \in H (u_i \text{ is defined})$
- Agreement $\forall i, j \in H (u_i = u_j)$
- Validity $\exists v (\forall i \in H (v_i = v)) \implies (\forall i \in H (u_i = v))$

Consensus Properties

- Termination $\forall i \in H (u_i \text{ is defined})$
- Agreement $\forall i, j \in H (u_i = u_j)$
- Validity $\exists v (\forall i \in H (v_i = v)) \implies (\forall i \in H (u_i = v))$
- Strong Validity $\forall i \in H \exists j \in H (u_i = v_j)$

Honest Majority is Necessary, I



Consider an adversary that performs one of the following with probability $\frac{1}{3}$

Honest Majority is Necessary, II

- If consensus protocol secure:
 - Adversary corrupts A_0 : output of honest parties (that belong to A_1) should be 1.
 - Adversary corrupts A_1 : output of honest parties (that belong to A_0) should be 0.
 - Adversary corrupts no-one: output of all parties should be the same.
- Adversary corrupts each set with prob. $\frac{1}{3}$ and instructs corrupted parties to follow the protocol
 - honest parties cannot distinguish between honest/corrupted parties
- If all parties output same value: validity is violated with prob. at least $\frac{1}{3}$
- If all parties output different value: consistency is violated with prob. at least $\frac{1}{3}$

Is Honest Majority Sufficient?

- Two important scenarios have been considered in the consensus literature.
 - Point to point channels. **No setup.**
 - Point to point channels. **With setup.**
- The setup provides a correlated private initialization string to each participant; *it is assumed* to be honestly produced.

Setup and Network

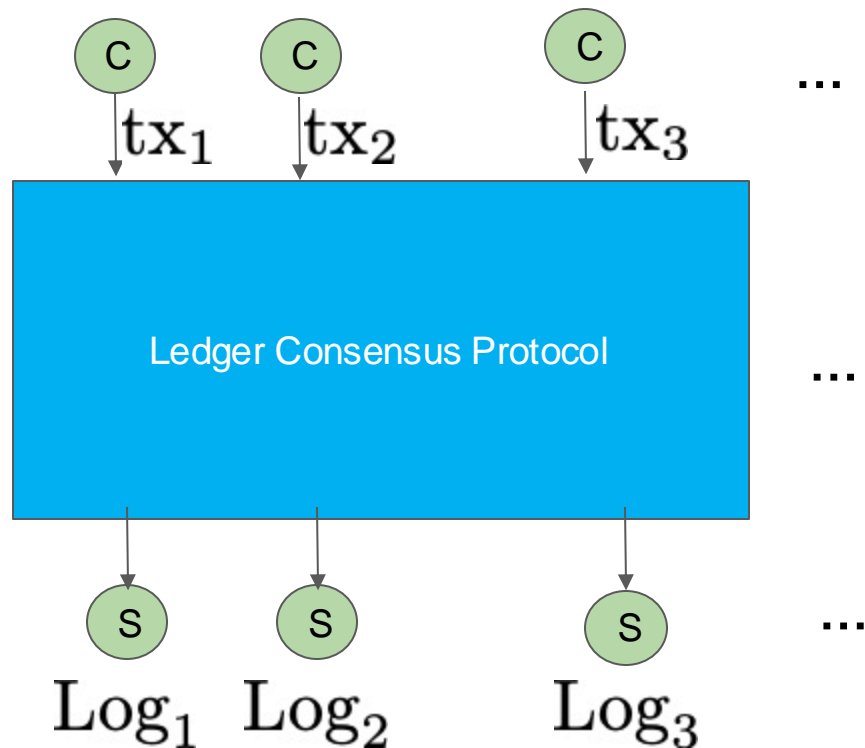
Setup/Network	Synchrony	Partial Synchrony
No Setup	$t < n/3$	$t < n/3$
With Setup	$t < n/2$	$t < n/3$

We know consensus can be achieved, assuming the above bounds on adversarial parties.

The typical setup and network configuration in classical consensus protocols

- Setup: a public-key directory
 - Parties have signing and verification keys for a digital signature scheme.
 - Each party knows every other party's verification key.
- Network: point-to-point channels
 - Synchronous, partially synchronous, or asynchronous

Ledger Consensus



Related to state machine
replication [Sch90]

Ledger Consensus Properties

- Consistency:

$$\forall i, j \in \mathbf{H}, t, t' : (\text{Log}_i[t] \not\preceq \text{Log}_j[t']) \rightarrow (\text{Log}_j[t'] \preceq \text{Log}_i[t])$$

- Liveness:

$$(\forall i \in \mathbf{H} : \text{tx} \in I_i[t]) \rightarrow (\forall i \in \mathbf{H} : \text{tx} \in \text{Log}_i[t + u])$$

$I_j[t] =$ Transaction Input of party j at time t consistent with its log

$\text{Log}_j[t] =$ Log of party j at time t

[.. and there are more properties of interest here.. e.g., fairness of transaction serialization, exporting a clock,...]

Solving ledger consensus

(solving SMR, actually)

Practical Byzantine Fault Tolerance

Miguel Castro and Barbara Liskov
Laboratory for Computer Science,
Massachusetts Institute of Technology,
545 Technology Square, Cambridge, MA 02139
`{castro,liskov}@lcs.mit.edu`

arbitrary behavior. Whereas previous algorithms assumed a synchronous system or were too slow to be used in practice, the algorithm described in this paper is practical: it works in asynchronous environments like the Internet and incorporates

- Its operation & threat model features:
 - Asynchronous communication, point-to-point channels.
 - Authenticated setting: parties are able to identify message sources
 - Fixed number of participants
 - Computationally bounded adversary in the sense it cannot forge digital signatures.
 - Setup assumption: PKI
 - Adversary controls less than $\frac{1}{3}$ of participants (required)
 - No privacy properties for servers
 - communication complexity $O(n^2)$ in best case, where n is the number of participants
 - No protocol incentives

Solving ledger consensus, II

Practical Byzantine Fault Tolerance

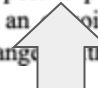
Miguel Castro and Barbara Liskov
*Laboratory for Computer Science,
Massachusetts Institute of Technology,
545 Technology Square, Cambridge, MA 02139*
{castro,liskov}@lcs.mit.edu

arbitrary behavior. Whereas previous algorithms assumed a synchronous system or were too slow to be used in practice, the algorithm described in this paper is practical: it works in asynchronous environments like the Internet and incorporates

Bitcoin: A Peer-to-Peer Electronic Cash System



Satoshi Nakamoto
satoshin@gmx.com
www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing



runs uninterruptedly over the Internet since 2009 assuming bounded delays

Bitcoin Blockchain

- A **ledger consensus** protocol, where
 - participants called “miners” run an intensive proof of work (PoW) algorithm at every step.
- Its threat model features:
 - **Bounded delay** communication by ‘diffusion’ of messages (no point-to-point channels)
 - No ability to **identify message** sources - unauthenticated setting
 - **Unknown** and **fluctuating number** of participants ... **dynamic availability** 
 - **Computationally bounded** adversary possessing $\sim < 50\%$ of resources
 - Setup: **public random string** (no PKI)
 - Tolerates potential **spikes** of adversarial majority... **self-heals** 
 - **Privacy for** maintainers... protocol transcript maintains anonymity/unlinkability of participants
 - **Optimal amortized communication complexity of single multicast per block.**
 - Protocol **ince**

[Nakamoto2009]

When bitcoin appeared the above threat model was an unexplored territory for consensus

The threat model was also not spelled out anywhere in the original works — protocol was given as is


The Bitcoin “backbone”

- The core of the bitcoin protocol
 - The chain validation predicate.
 - The chain selection rule (max-valid)
 - The proof of work function.
 - The main protocol loop
- Protocol is executed by “miners”

[GKL2015] Garay, Kiayias, Leonardos. The Bitcoin Backbone Protocol: Analysis and Applications.

Model

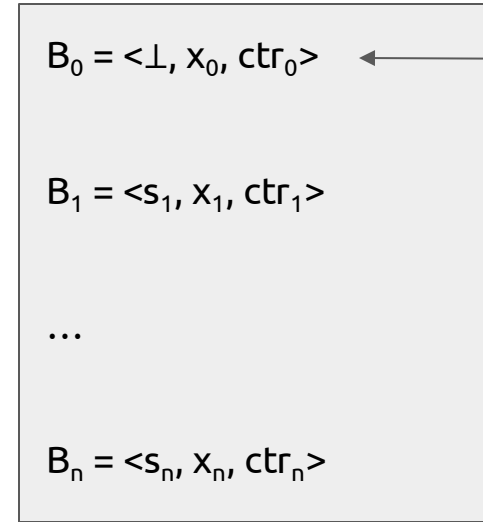
- Assume there are n parties running the protocol
- Synchronous
- Each party has a quota of q queries to the function $H(\cdot)$ in each round
- A number of t parties are controlled by an adversary (a malicious coalition)
 - Security arguments are *for any adversary*



ALGORITHM 3: The *PoW* function, parameterized by q, T, s_{init} and hash functions $H(\cdot), G(\cdot)$. The input is (x, C) .

```
1: function pow( $x, C$ )
2:   if  $C = \varepsilon$  then                                     ▷ Determine PoW instance
3:      $s \leftarrow s_{\text{init}}$                                    ▷ Genesis block initialization
4:   else
5:      $(\langle s', x', ctr' \rangle, s) \leftarrow \text{head}(C)$ 
6:   end if
7:    $ctr \leftarrow 1$ 
8:    $B \leftarrow \varepsilon$ 
9:    $h \leftarrow G(s, x)$ 
10:  while  $(ctr \leq q)$  do                                     ▷ This  $H(\cdot)$  invocation subject to the  $q$  bound
11:     $s_{\text{new}} \leftarrow H(ctr, h)$ 
12:    if  $(s_{\text{new}} < T)$  then
13:       $B \leftarrow \langle s, x, ctr \rangle$ 
14:      break
15:    end if
16:     $ctr \leftarrow ctr + 1$ 
17:  end while
18:  if  $B \neq \varepsilon$  then
19:     $C \leftarrow C || (B, s_{\text{new}})$                                ▷ Extend chain
20:  end if
21:  return  $C$ 
22: end function
```

Blockchain



genesis block

$$s_1 = H(\text{ctr}_0, G(x_0, \text{"label"})) = s_{\text{init}}$$

$$s_i = H(\text{ctr}_{i-1}, G(x_{i-1}, s_{i-1}))$$

$$X_C = \langle x_0, x_1, \dots, x_n \rangle$$


$$C^k = \langle B_0, B_1, \dots, B_{n-k} \rangle$$

ALGORITHM 1: The *chain validation predicate*, parameterized by q, T , the hash functions $G(\cdot), H(\cdot)$, and the *content validation predicate* $V(\cdot)$. The input is C .


```

1: function validate( $C$ )
2:    $b \leftarrow V(x_C)$ 
3:   if  $b \wedge (C \neq \varepsilon)$  then                                 $\triangleright$  The chain is non-empty and meaningful w.r.t.  $V(\cdot)$ 
4:      $(\langle s, x, ctr \rangle, s') \leftarrow \text{front}(C)$ 
5:      $s_{\text{front}} \leftarrow s'$ 
6:     repeat
7:       if  $\text{validblock}_q^T(\langle s, x, ctr \rangle) \wedge (H(ctr, G(s, x)) = s') \wedge (s_{\text{front}} = s')$  then
8:          $s_{\text{front}} \leftarrow s$                                  $\triangleright$  Retain hash value
9:          $C \leftarrow C^{\uparrow 1}$                                      $\triangleright$  Remove the head from  $C$ 
10:         $(\langle s, x, ctr \rangle, s') \leftarrow \text{head}(C)$ 
11:      else
12:         $b \leftarrow \text{False}$ 
13:      end if
14:    until  $(C = \varepsilon) \vee (b = \text{False})$ 
15:  end if
16:  return  $(b \wedge (s_{\text{front}} = s_{\text{init}}))$ 
17: end function

```



validblock predicate:



ALGORITHM 2: The function that finds the “best” chain, parameterized by function $\max(\cdot)$. The input is $\{C_1, \dots, C_k\}$.

```
1: function maxvalid( $C_1, \dots, C_k$ )
2:    $temp \leftarrow \varepsilon$ 
3:   for  $i = 1$  to  $k$  do
4:     if validate( $C_i$ ) then
5:        $temp \leftarrow \max(C_i, temp)$ 
6:     end if
7:   end for
8:   return  $temp$ 
9: end function
```

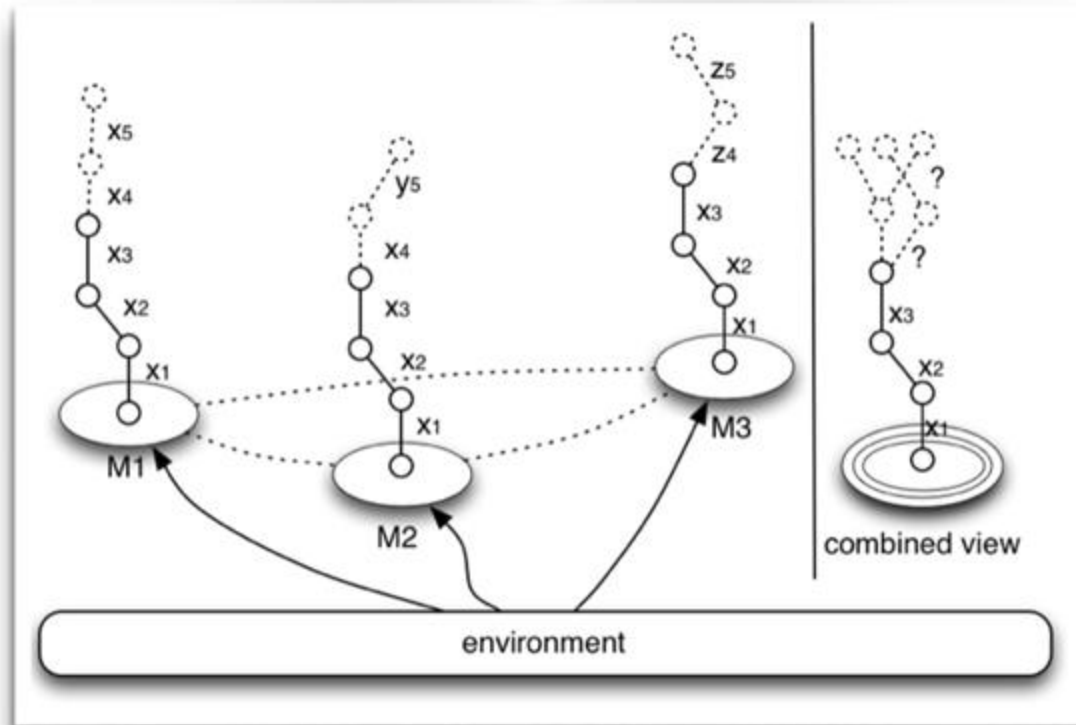
ALGORITHM 4: The Bitcoin backbone protocol's main loop, executed every round when a party is activated by the environment, parameterized by the *input contribution function* $I(\cdot)$ and the *chain reading function* $R(\cdot)$. At the onset it is assumed "init= True".

```
1: if (init) then
2:    $C \leftarrow \varepsilon$ 
3:    $st \leftarrow \varepsilon$ 
4:    $round \leftarrow 1$ 
5:    $init \leftarrow \text{False}$ 
6: else
7:    $\tilde{C} \leftarrow \text{maxvalid}(C, \text{any chain } C' \text{ found in RECEIVE})$ 
8:   if INPUT contains READ then
9:     write  $R(\tilde{C})$  to OUTPUT            $\triangleright$  Produce necessary output before the PoW stage.
10:  end if
11:   $\langle st, x \rangle \leftarrow I(st, \tilde{C}, round, \text{INPUT}, \text{RECEIVE})$             $\triangleright$  Determine the  $x$  value.
12:   $C_{\text{new}} \leftarrow \text{pow}(x, \tilde{C})$ 
13:  if  $C \neq C_{\text{new}}$  then
14:     $C \leftarrow C_{\text{new}}$ 
15:    DIFFUSE( $C$ )            $\triangleright$  Send the chain in case of adoption/extension.
16:  else
17:    DIFFUSE( $\perp$ )            $\triangleright$  Signals the end of the round to the diffuse functionality.
18:  end if
19:   $round \leftarrow round + 1$ 
20: end if
```

Basic Properties

- Common Prefix
- Chain Quality
- Chain Growth

Common Prefix, I



Common Prefix, II

(strong common prefix / consistency)

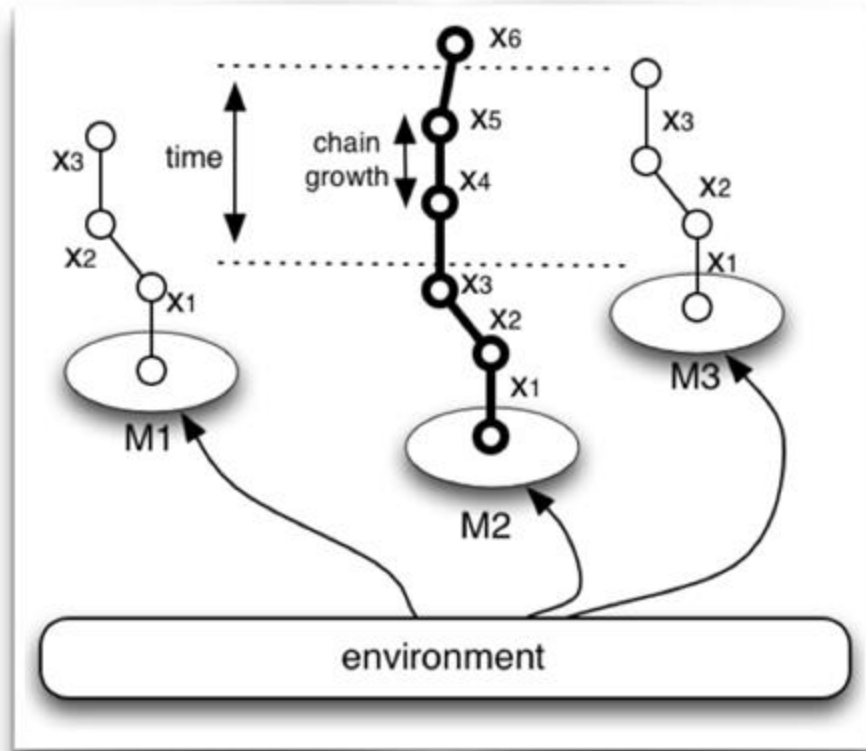
$$\forall r_1, r_2, (r_1 \leq r_2), P_1, P_2, \text{ with } \mathcal{C}_1, \mathcal{C}_2 : \mathcal{C}_1^{\lceil k} \preceq \mathcal{C}_2$$

- The property holds true, in a probabilistic sense, with an error that decays exponentially in k

Racing Attacks

- Attacker splits from the main chain and tries to overtake the “honest chain”
=> Common prefix breaks
- Intuition why the attack is a small probability event:
concentration bounds help honest parties

Chain Growth, I



Chain Growth, II

Parameters $\tau \in (0, 1), s \in \mathbb{N}$

In any period of s rounds at least τs blocks are added to the chain of an honest party P .

- The property holds true in a probabilistic sense with an error probability that exponentially decays in s

$\tau \approx$ probability at least one honest party finds a POW in a round

Abstention Attacks


- Attacker stops producing blocks

=> Chain growth stops

- Intuition why the attack is a small probability event:

honest parties will eventually issue blocks

Chain Quality, I



Chain Quality, II

Parameters $\mu \in \{0, 1\}, \ell \in \mathbb{N}$

The ratio of blocks of an ℓ -long segment of an honest chain produced by the adversary is bounded by $(1 - \mu)\ell$

- The property holds true probabilistically with an error that exponentially decays in ℓ

$$\mu \approx \frac{n - 2t}{n - t}$$

Block Withholding Attacks

- Attacker mines privately and releases their block at the same time an honest party releases its own block
- Assuming honest propagation favours the adversary, the honest block is dropped, reducing chain quality
- Intuition why the attack is a small probability event:
over time the adversary cannot produce blocks at the same rate as honest parties (to compete with them)

Establishing a Ledger Consensus from a Blockchain

- Consistency \leftarrow *(strong) Common Prefix*
 - need to exclude k most recent blocks
- Liveness \leftarrow *Chain Growth* and *Chain Quality*
 - leave sufficient time for chain to grow
 - apply chain quality to ensure that at least one honest block is included

Ledger Consensus vs. Consensus

- What is the connection?
 - ledger is an ever-going protocol with inputs (e.g., transactions) continuously coming from also external sources
 - consensus is a one-shot execution
- Is it possible to reduce consensus to the ledger? Is it possible to reduce the ledger to consensus?
 - (See the [GKL paper](#) for more details)

Hash operations

- Consider a regular PC (30 MHash / sec)
- With expectation of 2^{78} hashing operations (current difficulty – 10/24), mining a block will require ~ 320 *million* years.


https://en.bitcoin.it/wiki/Non-specialized_hardware_comparison

<https://www.coinwarz.com/mining/bitcoin/difficulty-chart>

Parallelising mining

- Bitcoin's Proof of Work can be parallelized
- Parties tend to form mining *pools*
 - Instead of working separately, work **together** to solve PoW for the same block.
 - By collecting “**shares**” (small hashes of the block that are not quite as small as needed) one can prove how much they contributed.

Bitcoin mining pools





Recall: PoW algorithm

```
int counter;  
counter = 0  
while Hash(data, counter) > Target  
    increment counter  
return counter
```


Dynamic Availability

- So far: n nodes maintain the blockchain
- This number may change over time:
 - new users enter the system
 - existing users leave
- The change over time can be **dramatic**
- The Bitcoin blockchain handles this, by adjusting the target (difficulty) of the Proof of Work algorithm

Target difficulty / Total hash rate over time



Adjusting the difficulty

“maxvalid” rule is changed

s.t. parties adopt **chain with highest difficulty** linearly related to:

$$\sum_i \frac{1}{T_i}$$

The f parameter [GKL15]

f = probability of producing a block in a round of interaction

- f depends on:
 - target T
 - number of miners
 - duration of round
- If f becomes too small, parties do not progress
 - Chain growth slows
 - Liveness is hurt
- If f becomes too large, parties “collide” often
 - Attacker can exploit network scheduling of message delivery to create forks
 - Consistency is hurt
- To resolve this dynamically, Bitcoin **recalculates** T to keep f constant

Target recalculation

$$\text{next target} = \begin{cases} \frac{1}{\tau} \cdot T & \text{if } \frac{n_0}{n} \cdot T_0 < \frac{1}{\tau} \cdot T; \\ \tau \cdot T & \text{if } \frac{n_0}{n} \cdot T_0 > \tau \cdot T; \\ \frac{n_0}{n} \cdot T_0 & \text{otherwise} \end{cases}$$

- Recalculation occurs at the end of every “epoch”
 - m : epoch length in blocks (in Bitcoin: 2016)
- n_0 : estimation of number of ready parties at the system’s onset (party=CPU)
- T_0 : initial target
- τ : recalculation threshold parameter (in Bitcoin: 4)
- T : target in effect

- $n = m/(pTM)$: the “effective” number of parties in the epoch
 - M : last epoch’s duration based on block timestamps
 - pT : probability of a single party being successful in PoW in a round

Clay pigeon shooting game



Clay pigeon shooting game

- Suppose you shoot on targets successively against an opponent
 - your success probability: $p = 0.3$
 - your opponent's success probability: $q = 0.4$
 - you shoot in sequence 1000 targets
 - winner is the one that got the most hits
- What is your probability of winning?

Analysis, I

- Consider the random variable $Z_i = X_i + X_{i-1} + \dots + X_1$
 - $X_i=1$ you get ahead, probability $p(1-q)$
 - $X_i=-1$ opponent gets ahead, probability $q(1-p)$
 - $X_i=0$ draw, probability $pq + (1-p)(1-q)$
- $E[X_i] = p(1-q) - q(1-p) = p - q < 0$
- $E[Z_n] = n(p - q)$
- Hoeffding inequality states: $\Pr[Z_n \geq \varepsilon + n(p - q)] \leq \exp(-\varepsilon^2 / 2Bn)$, for $\varepsilon > 0$ and $-B \leq X_i \leq B$
- Set $\varepsilon = n(q - p)$, $B=1$, and we obtain $\Pr[Z_n \geq 0] \leq \exp(-(q-p)^2 n / 2)$
- For our numbers ($q=0.4$, $p=0.3$, $n=1000$), $\Pr[Z_n \geq 0] \leq 0.007 \leq 1\%$

Analysis, II

- Now you are given a choice:
 - decrease the size of the clay pigeon target by a ratio β
 - augment your “kills” by multiplying with $1/\beta$
 - your accuracy is linear with β
 - your opponent will keep playing in the same way as before
- Do you prefer to play like this or it makes no difference?

Analysis, III

- Consider the random variable $Z_i = X_i + X_{i-1} + \dots + X_1$
 - $X_i=1/\beta$ you get ahead, probability $\beta p(1-q)$
 - $X_i=-1$ opponent gets ahead, probability $q(1-\beta p)$
 - $X_i=1/\beta-1$ both win, probability $\beta p q$, and $X_i=0$ both lose, probability $(1-\beta p)(1-q)$
- $E[X_i] = (1/\beta)\beta p(1-q) - q(1-\beta p) + (1/\beta-1)\beta p q = p - q < 0$
- $E[Z_n] = n(p - q)$
- Hoeffding inequality states: $\Pr[Z_n \geq \varepsilon + n(p - q)] \leq \exp(-\varepsilon^2 / 2Bn)$, for $\varepsilon > 0$ and $-B \leq X_i \leq B$
- Set $\varepsilon = n(q - p)$, $B=1/\beta$, and we obtain $\Pr[Z_n \geq 0] \leq \exp(-\beta(q-p)^2 n / 2)$
- Observe: as $\beta \rightarrow 0$, tail bound deteriorates to ≤ 1 .

The Difficulty Raising Attack

- The recalculation threshold (τ) is essential
- Without it, an adversary that has a minority of hashing power:
 - Creates a private, artificially difficult chain
 - Similar to clay pigeon shooting game, this increases the variance in its block production rate
 - Overcoming the chain of the honest parties becomes a non-negligible event

[B13] Lear Bahack. Theoretical Bitcoin Attacks with less than Half of the Computational Power (draft)